

Computer Security

Coursework Exercise CW3

Software Security

The goal of this coursework is to gain practical experience with attacks that exploit software vulnerabilities, in particular buffer overflows. A buffer overflow occurs when a program attempts to write data beyond the bound of a fixed-length buffer – this can be exploited to alter the flow of the program, and even execute arbitrary code.

You are given five exploitable programs, which are installed within the virtual machine provided for this coursework. Each of you will work on slightly different exploitable programs. You will find the five programs you need to work on in the `/task{1-4}/[yourUUN]` and `/task5` directories on the provided virtual machine, with `setuid` for a respective user set.¹ In each case, a secret should be extracted by exploiting the program in some way, by supplying it maliciously crafted inputs, either to spawn a shell, or extract the secret directly.

1 Virtual Machine Layout and Setup

For this assignment we are using a Virtual Machine called `binary-exploits` (how original!) inside VirtualBox. The machine comes as an `*.ova` file, which includes the virtual hard drive as well as all necessary settings for VirtualBox. To get the coursework running:

1. Download VirtualBox (version 5.2.42) at <https://www.virtualbox.org/>
2. Download the virtual appliance at <https://edin.ac/3lr5Euv>
3. Once VirtualBox is running, select File ↔ Import Appliance and select the file `binary-exploits.ova`.

For `binary-exploits` your username and password are both `student`, and you have root access via `sudo`. Be cautious with how you use it, however, as you will not have root access on our automarker!

In order to allow these attacks to be practically and reproducibly executed, address space layout randomisation (ASLR) has been disabled in the VM, ensuring that in each execution, the same parts of a program get assigned to a consistent memory location.

To further protect against buffer overflow attacks and other attacks that use shell programs, many shell programs automatically drop their privileges when invoked. Therefore, even if you can “fool” a privileged `setuid` program to invoke a shell, you might not be able to retain the privileges within the shell. This protection scheme is implemented in `/bin/bash`. In Ubuntu, `/bin/sh` is actually a symbolic link to `/bin/bash`. For this coursework, and in order to really realise how much damage this sort of attacks can do, we use another shell program (the `zsh`), instead of `/bin/bash`. The preconfigured Ubuntu virtual machines contains a `zsh` installation.

Templates On first starting the VM, templates for each of the exercise submissions should be found in the `cw3/` folder in your home directory. For each task you will find a script `task{1-5}.sh`. This script will be executed by our automarker, and should output the secret to its `stdout`. Beyond this, you are free to use what you want – provided the tools are preinstalled on the VM. This includes using your own compiled binaries if you wish, although you will not get partial marks if we can’t see the source! Notable tools installed include netcat (`nc`), and `gdb`.

¹Note that for task 5 you are all working on the same program in `/task5`.

Important You must run your attacks inside the provided VM. It is very important that your attack programs work in the provided **binary-exploits** VM. This is because, the compiler version, the operating system and the installed libraries will affect the exact location of code on the stack. The VM provides an identical environment to the one in which we will test your code for marking it.

Some of the template scripts include a part using the **env** command to clear the environment. This is to ensure the environment our automarker runs the attack in, and the environment you test it in are the same. Take care to also make this consistent with your gdb environment! If you remove this part, you do so at your own peril, and you may experience an attack that works fine for you being turned down by our automarker.

2 Shellcode

In task 3 you will have to deploy some shellcode, and while we do not ask you to come up with this yourself, we will briefly elaborate what this code does. These are x86 instructions, that, when executed, have the same affect as executing the following C function:

```
int main() {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    setreuid(geteuid(), geteuid());
    execve(name[0], name, NULL);
}
```

This program does two things: First, it sets the current user id the the effective user ID. This ensures that the shell does not drop privileges when it is started, i.e. reverting to an unprivileged user. In particular bash does this, which supplies **/bin/sh** in our VM. Second, it executes **/bin/sh** directly.

Good shellcode does not contain NULL bytes, as these will not typically be copied by C programs – for this reason just compiling the above is usually not sufficient. We provide the below shellcode, and in task 3, this is already available in the task template.

```
/* geteuid() */
"\x6a\x31" // push $0x31;
"\x58" // pop %eax;
"\x99" // cltd;
"\xcd\x80" // int $0x80;

/* setreuid() */
"\x89\xc3" // mov %eax, %ebx;
"\x89\xc1" // mov %eax, %ecx;
"\x6a\x46" // push $0x46;
"\x58" // pop %eax;
"\xcd\x80" // int $0x80;

/* execve("/bin/sh", 0, 0) */
"\xb0\x0b" // mov $0x0b, %al;
"\x52" // push %edx;
"\x68n/sh" // push $0x68732f6e;
"\x68//bi" // push $0x68732f6e;
"\x89\xe3" // mov %esp, %ebx;
"\x89\xd1" // mov %edx, %ecx;
"\xcd\x80" // int $0x80;
```

3 The Vulnerable Programs

Task 1 This is just a warm-up. The program in `/task1/[yourUUN]/vuln` waits for a password (stored in `/task1/password.txt`), and if the correct password is entered by the user, it prints the secret. The program has a buffer overflow vulnerability, making it possible – without knowing the password – to allow it to log in. **Important:** the real password will be changed for marking! The `pidgeon` and `rooster` various **will not**.

Task 2 The program in `/task2/[yourUUN]/vuln` takes the user's name as a command line argument, copies it to a buffer, and then welcomes the user. This program is vulnerable to a buffer overflow attack. Your attack script should call this program with a carefully crafted argument, such that it overwrites the return address on the stack, and returns to the `read_secret` function, instead of back to `main`. Stack canaries are not enabled in this task.

Task 3 For the program in `/task3/[yourUUN]/vuln`, there is no helpful function for extracting information, meaning you must inject your own shellcode, jump to this on-stack instead of a pre-existing function. Furthermore, from this task on, the program is compiled with GCC's stack protection enabled, which inserts stack canaries after the return address on stack. Your attack script will input a malicious argument, such that the program loads and jumps into your shellcode, while avoiding these canaries. The template provided already contains some boilerplate code to then read the secret from this shell.

Task 4 The program in `/task4/[yourUUN]/vuln` functions similarly to the previous one, however execution of code on the stack has been disabled! Instead, you should perform a `return-to-libc` attack. In particular, you should overwrite the stack such that execution returns into the `libc` function `system()`, and make this function believe it received `"/bin/sh"` as an argument. For this purpose, you'll need to find the locations of these in memory, and analyse how calls to `system()` would ordinarily function. Your program should use this return technique to spawn a shell – in this case the program itself forces that the real uid is set, ensuring the shell will be as the `root` user.

Task 5 This task does not have a binary to directly exploit, but instead runs a small TCP server on the local port 4040. This program is found in `/task5/vuln.py`. It accepts requests and read sections of a secret file, but denies requests to some parts. It is however vulnerable to an attack allowing you to read this part too! Your attack script should connect to the server using netcat, and submit a malicious request, extracting the secret string.

4 Submission Instructions

Once you have completed your exploits, you should first test them against a clean `binary-exploits` VM. To retrieve your attack scripts `task{1-5}` from the VM you just need to change the settings Devices \leftrightarrow Shared Folders \leftrightarrow Shared Folders Settings. If you are satisfied, and want to submit, you go on Learn. In the Assessment content area, you will find in the Assignment Submission folder, the Coursework 3 - Software security submission box. Learn does not support `.sh` files. You will need to create a zip archive with your attack scripts (create an archive of your shell script files rather than an directory), and you will name it `UUN.zip` which you will be able to submit. You will find further guidance on how to submit here:

<https://edin.ac/3rZK2b1>

The submission deadline is ~~1st April, at 16:00~~ 5th April, at 16:00. **Please be aware:** The actual marking will generate new random secrets, different for every student. Your secrets must therefore come from executing the attack on your exploitable programs specifically, and not, for instance, reading them with root and simply echoing them!

A Running on DICE Machines

Note - The setup given here is adapted from a previous iteration of a similar assignment. It is in theory

usable (it worked for me), but it might require minor tweaks, that we'll have to discover together :-)

The VM is accessible from all DICE machines here `/disk/scratch/software_security/`, in a compressed virtual appliance format. VirtualBox is already installed on all DICE computers. Once VirtualBox is running, select File \leftrightarrow Import Appliance and select the file `/disk/scratch/software_security/binary-exploits.ova`. The only property in the configuration screen (Appliance settings) you should change is the location of the virtual hard drive. On DICE computers a home directory account does not have enough quota to store the hard drive, so you will need to store the virtual disk locally on the workstation. The disk takes up about 8G once expanded. In the Appliance settings screen change the path from, e.g.

`/afs/inf.ed.ac.uk/user/sXX/sXXXXXX/VirtualBoxVMs/`

to

`/tmp/sXXXXXX/VirtualBoxVMs/`

(you may need to create the directory `/tmp/sXXXXXX` first). The import operation can take several minutes and need about 8.5G of disk space per machine.

On DICE machines the virtual disk has to be stored on the local disk. That means that you will not be able to access the virtual machine on any other DICE computer than the one you did this setup on. If you want to use your virtual machine from another computer you will have to log in to the machine using `ssh -X <computername>`, or copy the disk files to another machine with `scp`; so make sure you remember the name of the computer you used!